

Lawrence Lin

925-918-2478 | lawrencedlin@gmail.com | [linkedin.com/in/llinnn](https://www.linkedin.com/in/llinnn) | github.com/lawrencedlin

EDUCATION

University of San Francisco

San Francisco, CA

M.S. Data Science

July 2021 - August 2022 (Expected)

Courses: Advanced Machine Learning, Distributed Computing, Distributed Data Systems, Time Series

University of California, Santa Barbara

Santa Barbara, CA

B.S. Statistics

August 2017 - June 2021

Courses: Machine Learning, Bayesian Statistics, Data Structures and Algorithms, Stochastic Processes

EXPERIENCE

Data Science Intern

November 2021 - Present

Walmart Labs

Sunnyvale, CA

- Analyzed millions of customer's seasonal purchase behavior for festivals using Apache Spark and Seaborn
- Worked on feature engineering and data cleaning including categorical encoding and sequence creation
- Developed a Transformer Neural Network Model in TensorFlow to sequentially recommend top-k items
- Performed hyper-parameter grid-search by deploying model on a Google Dataproc cluster and achieved an AUC of 0.88, Mean Reciprocal Rank of 0.56, and Normalized Discounted Cumulative Gain of *INSERT BELIEVABLE NUMBER HERE*
- Evaluated user, item, and time embedding quality by analyzing k-closest embedding by Euclidean distance

Research Assistant

January 2021 - June 2021

Sansum Diabetes Research Institute

Santa Barbara, CA

- Visualized diabetes severity by zip code using GeoPandas and Folium
- Tested for statistically significant differences in HbA1c among demographic groups using ANOVA and Welch's t-test with Bonferonni Correction
- Modeled HbA1c with LASSO and OLS regression models achieving an R^2 of 0.77
- Authored weekly written reports and presented insights to SDRI researchers

PROJECTS

Implicit Rating Prediction | *Pytorch, FastAI*

- Developed a Matrix Factorization model in PyTorch and a Tabular Neural Network model in FastAI
- Trained models with negative sampling algorithms and cyclical learning rates
- Achieved 1st place on Kaggle leaderboard with a binary cross-entropy loss of 0.4032

Twitter and Reddit Sentiment Analysis | *AWS, Databricks, Spark, MongoDB, BERT*

- Extracted over a year of reddit comments and tweets mentioning an controversial celebrity using REST APIs and stored data in Amazon S3 and a MongoDB cluster
- Created new features from text using pre-trained BERT emotion and sentiment models from HuggingFace
- Predicted YouTube weekly viewership on engineered sentiment and emotion features using Random Forest and Gradient-Boosted Regression models through SparkML on Databricks cluster

Feature Importance | *Scikit-Learn, NumPy, Pandas*

- Implemented Spearman correlation, Principal Components Analysis, Permutation and Drop-column importance
- Visualized the cross-validated R^2 of a Gradient-Boosted Regressor trained on k most important features
- Implemented automatic forward feature selection algorithm using permutation importance
- Calculated variance and empirical p-value of feature importances using bootstrap samples

TECHNICAL SKILLS

Languages: Python, R, C++, SQL (Postgres), NoSQL (Mongo) HTML/CSS, Bash

Frameworks: Hadoop Ecosystem (HDFS, YARN, Spark, SparkMLib, HiveQL) TensorFlow, PyTorch, FastAI, Scikit-Learn, Statsmodels, Scipy, Numpy, Pandas, Matplotlib, Seaborn, Flask, BeautifulSoup, Selenium, H2o, HuggingFace

Developer Tools: Git, Docker, Google Cloud Platform, Amazon Web Services, DataBricks