

Lawrence Lin

San Francisco, CA | lawrencedlin@gmail.com | [linkedin.com/in/llinnn](https://www.linkedin.com/in/llinnn) | github.com/lawrencedlin

EDUCATION

University of San Francisco

July 2021 - July 2022

M.S. Data Science

Courses: Advanced Machine Learning, Deep Learning, Relational Databases, Time Series Analysis, A/B Testing

University of California, Santa Barbara

August 2017 - June 2021

B.S. Statistics

Courses: Machine Learning, Bayesian Statistics, Stochastic Processes, Data Structures and Algorithms

EXPERIENCE

Data Scientist

July 2022 - Present

Walmart Global Tech

Sunnyvale, CA

- Built from scratch a Walmart+ member feature pipeline driven by Spark and orchestrated by cron
- Trained a Gradient-Boosted Trees churn propensity model using Walmart+ member feature platform
- Powered intelligent order quality checks on high churn-risk members with churn model causing positive lift in GMV and re-orders

Data Science Intern

November 2021 - July 2022

Walmart Global Tech

Sunnyvale, CA

- Discovered peak festival shopping activity windows for millions of customers' using clustering algorithms
- Independently performed feature engineering and data cleaning on distributed datasets
- Developed and trained a Transformer Neural Network Model in TensorFlow to make personalized season-aware recommendations using historical purchases with an AUC of 0.88
- Validated time embedding quality by finding high average cosine similarities over 7-day windows

Research Assistant

January 2021 - June 2021

Sansum Diabetes Research Institute

Santa Barbara, CA

- Visualized Californian zip codes most severely impacted by diabetes using GeoPandas and Folium
- Tested for statistically significant differences in blood sugar levels among Hispanic population using ANOVA
- Modeled blood sugar levels with LASSO and OLS regression models achieving an R^2 of 0.77

PROJECTS

Implicit Rating Prediction | *Pytorch, FastAI*

- Developed a Matrix Factorization model and a Tabular Neural Network model to predict implicit hotel ratings
- Achieved 1st place on Kaggle leaderboard with a binary cross-entropy loss of 0.4032

Twitter and Reddit Sentiment Analysis | *AWS, Databricks, Spark, MongoDB, BERT*

- Scraped over a year of reddit comments and tweets and stored data in Amazon S3 and a MongoDB cluster
- Engineered new features from social media with BERT emotion and sentiment models from Hugging Face
- Predicted YouTube weekly viewership on engineered sentiment and emotion features using Random Forest and Gradient-Boosted Regression models through SparkML on Databricks cluster

SKILLS

Languages: Python, R, C++, SQL (Postgres), NoSQL (Mongo) HTML/CSS, Bash

Frameworks: Hadoop Ecosystem (HDFS, YARN, Spark, SparkMLib, HiveQL) TensorFlow, PyTorch, Scikit-Learn

Developer Tools: Git, Docker, Google Cloud Platform, Amazon Web Services, Databricks